

# etl\_script

March 31, 2025

import lib

```
[1]: import warnings
import pandas as pd
import sqlite3

warnings.filterwarnings('ignore')
```

## 1 Extract

```
[2]: # File paths
CUSTOMERS_CSV = "data/sources/customers.csv"
TRANSACTIONS_CSV = "data/sources/transactions.csv"
PRODUCTS_CSV = "data/sources/products.csv"
DB_NAME = "retail_data.db"

# Extract
customers = pd.read_csv(CUSTOMERS_CSV)
transactions = pd.read_csv(TRANSACTIONS_CSV)
products = pd.read_csv(PRODUCTS_CSV)
```

## 2 Transform

For each data source, the cleaned data frame will set the name `proc_*`

### 2.1 Customers

```
[3]: customers["email"].tail()

995          @example.com
996  anneparks@example.org
997    cjones@example.net
998  coreyhill@example.org
999    remember@example
Name: email, dtype: object
```

reference **gmail**'s mail naming convention

A valid email is a set of characters in format `username @ domain` whereas:

- a valid `username` which is:
  - start character must be a letter (a-z) or digit (0-9)
  - other character must be a letter (a-z), digit (0-9) or dot (.)
- a valid `domain` which is:
  - start with group of characters only letter (a-z) or digit (0-9)
  - next character is dot (.)
  - end with group of characters only letter (a-z) or digit (0-9)

row index 995 and 999 at above dataframe has invalid email address

## Processing

```
[4]: # valid email pattern
pattern = r"^[a-zA-Z0-9][a-zA-Z0-9.]+@[a-zA-Z0-9]+\.[a-zA-Z0-9]+$"
valid_emails = customers["email"].str.contains(pattern)
# customers with invalid email
invalid_email_customers = customers[~valid_emails]
# check invalid email
invalid_email_customers["email"].head(5)
```

```
[4]: 3      invalid_email@
    14      email.example.com
    17      machine@example
    21      email.example.com
    26      agreement@example
    Name: email, dtype: object
```

## Processed Customer

```
[5]: # customers with valid email
proc_customers = customers[valid_emails]
```

## 2.2 Product

```
[6]: products["category"].unique()
```

```
[6]: array(['fashion', 'electronics', 'Electronics', 'home', 'Home', 'Fashion'],
      dtype=object)
```

For this data set, we only need to lowercase to make them consistent.

## Processed Product

```
[7]: proc_products = products.copy()
proc_products["category"] = products["category"].apply(lambda x: x.lower())
# check
proc_products["category"].unique()
```

```
[7]: array(['fashion', 'electronics', 'home'], dtype=object)
```

## 2.3 Transaction

### 2.3.1 Deduplicate

```
[8]: transactions["transaction_id"].duplicated().sum()
```

```
[8]: 1000
```

```
[9]: transactions.duplicated().sum()
```

```
[9]: 1000
```

- If we got duplication on whole row contents -> keep only one unique row
- if duplication on transaction\_id -> source data has problems

In our situation, drop\_duplicates() is enough

```
[10]: dedup_transactions = transactions.drop_duplicates()
```

### 2.3.2 handle invalid transaction dates

```
[11]: # count how many date are in the format DD/MM/YYYY
transactions_dd_mm_yyyy =
    ↳ dedup_transactions[dedup_transactions["transaction_date"].str.
    ↳ contains(r"^\d{2}/\d{2}/\d{4}$")]
print("num transaction in format dd/mm/yyyy is: ", transactions_dd_mm_yyyy.
    ↳ count().max())
```

num transaction in format dd/mm/yyyy is: 159872

```
[12]: # count how many date are in the format YYYY-MM-DD
transactions_yyyy_mm_dd =
    ↳ dedup_transactions[dedup_transactions["transaction_date"].str.
    ↳ contains(r"^\d{4}-\d{2}-\d{2}$")]
print("num transaction in format yyyy-mm-dd is: ", transactions_yyyy_mm_dd.
    ↳ count().max())
```

num transaction in format yyyy-mm-dd is: 32040

```
[13]: # locate the other date format that is not in the two above format
invalid_transactions =
    ↳ dedup_transactions[~dedup_transactions["transaction_date"].str.
    ↳ contains(r"^\d{2}/\d{2}/\d{4}$|^\d{4}-\d{2}-\d{2}$")]
print("invalid_transactions count: ", invalid_transactions.count().max())
```

invalid\_transactions count: 8088

```
[14]: invalid_transactions['transaction_date']
```

```
[14]: 8          enough
      20        million
```

```

23         involve
27         single
29         first
...
199931     parent
199944  production
199963         most
199980     medical
199992         such
Name: transaction_date, Length: 8088, dtype: object

```

Most of them are text, we can skip it

We don't know for the date that in slash-break format is in dd/mm/yyyy or mm/dd/yyyy, try using pandas convert datetime function helper, if no `ValueError` exception is raised, our assumption is correct

```
[15]: _ = pd.to_datetime(transactions_dd_mm_yyyy["transaction_date"], format="%d/%m/%Y")
```

No error! So the slash-break datetime format is dd/mm/yyyy format.

Move forward to have a view on yyyy-mm-dd(or yyyy-dd-mm) format

```
[16]: transactions_yyyy_mm_dd["transaction_date"].head()
```

```

[16]: 3      2023-02-30
      4      2025-01-40
      5      2023-02-30
      13     2024-13-45
      21     2024-13-45
Name: transaction_date, dtype: object

```

```

[17]: # try to convert the date to datetime, if it fails, return these rows to na
      # then compare with original num rows
      # try %Y-%m-%d
      converted_transaction_yyyy_mm_dd = pd.
      # to_datetime(transactions_yyyy_mm_dd["transaction_date"], format="%Y-%m-%d",
      # errors="coerce")
      print(converted_transaction_yyyy_mm_dd.isna().count() ==
      # transactions_yyyy_mm_dd["transaction_date"].count())
      # try %Y-%d-%m
      converted_transaction_yyyy_dd_mm = pd.
      # to_datetime(transactions_yyyy_mm_dd["transaction_date"], format="%Y-%d-%m",
      # errors="coerce")
      print(converted_transaction_yyyy_dd_mm.isna().count() ==
      # transactions_yyyy_mm_dd["transaction_date"].count())

```

True

True

So we know that all of these %Y-%m-%d transaction date is invalid format

Final processed transactions:

```
[18]: # add these invalid transaction date to invalid_transactions df
invalid_transactions = pd.concat([invalid_transactions,
    ↪ transactions_yyyy_mm_dd])
# valid transactions
proc_transactions = dedup_transactions[~dedup_transactions["transaction_id"].
    ↪ isin(invalid_transactions["transaction_id"])]
proc_transactions["transaction_date"] = pd.
    ↪ to_datetime(proc_transactions["transaction_date"], format="%d/%m/%Y") #
    ↪ convert to datetime
```

### 3 Load

After processing data, From 3 original dataframe, we got proc\_\* dataframe which is cleaned data.

```
[19]: # Load the cleaned data into a SQLite database with three tables: customers,
    ↪ transactions, and products.
conn = sqlite3.connect(DB_NAME)
# write to db
proc_customers.to_sql("customers", conn, if_exists="replace", index=False)
proc_transactions.to_sql("transactions", conn, if_exists="replace", index=False)
proc_products.to_sql("products", conn, if_exists="replace", index=False)
```

```
[19]: 50000
```

```
[20]: pd.read_sql(
    sql="""
    SELECT sql
    FROM sqlite_master
    WHERE type = 'table' AND name = 'transactions'
    """, con=conn).values[0][0]
```

```
[20]: 'CREATE TABLE "transactions" (\n"transaction_id" TEXT,\n "customer_id" TEXT,\n
"transaction_date" TIMESTAMP,\n "amount" REAL\n)'
```

### 4 Data Aggregation

Write a SQL query to calculate the total transaction amount per customer and save the results in a new table called customer\_revenue.

```
[21]: customer_revenue = pd.read_sql(
    sql="""
    SELECT customer_id, sum(amount) as total_transaction_amount
    FROM transactions
    GROUP BY customer_id
```

```

        """ , con=conn)

customer_revenue.to_sql("customer_revenue", conn, if_exists="replace",
        ↪index=False)

```

[21]: 1000

```

[22]: pd.read_sql(
        sql="select * from customer_revenue", con=conn).head()

```

```

[22]:
           customer_id  total_transaction_amount
0  004454d5-8b28-4675-ac13-ffef982bc471      72340.74
1  009d8fce-6cef-4fb3-b7c3-59ce09c1e58a      73812.20
2  0131b090-2ca1-48f5-bf15-35b3f1923bdd      86753.93
3  014c1663-714f-4ea4-9d3a-9d477224044e      76594.32
4  01bd6b97-54eb-4413-85d8-c78fadf1e6f6      75030.63

```